THE COUNCIL METHOD

A field guide for getting real signal out of an AI conversation, instead of a generic answer.

Most people either paste a demand at an AI and take whatever comes back, or write an over-engineered mega-prompt and still get something generic. This is neither. It's four different situations, and the one specific move that actually works for each. Pick the one you're in right now.

A. I’ve got a messy idea and need to shape it into something real

B. I need something built or decided correctly, first try

C. I’ve already got something and need it properly stress-tested

D. I just want to think out loud and see what’s actually there

SHAPING A MESSY IDEA

A. Something in your head, not on paper yet

Don’t clean it up first. Say the actual messy version — the one in your head, not the one you’d put in an email. Structure is something you build together after, not something you owe it going in.

When it asks a clarifying question, that’s not friction — that’s the actual work starting. Answer it. “I don’t know, that’s what we’re figuring out” is a legitimate answer. It tells it exactly where you’re stuck.

The tell: if the answer you get back feels like it could’ve been written before you said anything specific, you haven’t given it enough yet. Go back and hand it something only you would know.

<< Back to start

GETTING IT RIGHT THE FIRST TIME

B. Constraints before solutions

Ban solutions until the constraints exist. Most people ask for the output first and correct it after. Flip it: spend the first stretch of the conversation only naming what's true — budget, audience, format, what "good" actually looks like — and refuse to let it, or yourself, jump ahead.

The tell you're doing it right: you're saying "not yet, tell me X first" more often than "no, try again."

When it gets something wrong because a constraint was missing, that's not the tool failing. That's a constraint you hadn't said out loud yet. Name it and move forward — don't relitigate the attempt.

[<< Back to start]

STRESS-TESTING SOMETHING REAL

C. Make it find the crack, not give an opinion

Assign it an explicit adversarial job, not a validating one. “Find what’s wrong with this” gets a completely different answer than “what do you think.”

Make it separate fact from convention. Most real errors hide in the gap between “this is true” and “this is just what we decided to do” — ask directly which one you’re looking at.

Don’t accept agreement as proof. If it agrees fast, push harder before you trust it — fast agreement is a property of the conversation, not of being correct.

Ask for the specific failure, not the general opinion. “Where exactly does this break” beats “does this look okay.”

<< Back to start

THINKING OUT LOUD, ON PURPOSE

D. Say the dumb version first

Say the dumb version first, out loud, before you've decided if it's worth saying. Most good material dies in the half-second before you type it, not after.

Its job here isn't to answer. It's to catch the one thing worth keeping in what you just said, and hand it back sharper than you gave it.

The tell: if you're slightly embarrassed by what you just said, you probably just said something real.

<< Back to start